

## **CS342 Fall 2016 – Project 5**

### **Exploring the FAT32 File System**

doc version 1.1

**Assigned:** Dec 15, 2016, Thu

**Due date:** Dec 29, 2016, Thu, 23:55

*Objectives: Learning FAT32 file system in detail. Understanding on-disk file system structures. Practicing formatting and mounting a file system, and parsing on-disk file system structures.*

In this project you will explore the FAT32 file system and write a program that will read and manipulate a FAT32 volume directly in raw mode. The project will be done in Linux and using C programming language. You can work in groups of 2 students each.

First read and learn about FAT32 file system and its on-disk structures: its volume information structure, directory record structure, FAT structure, etc.. You can find a lot of documentation about FAT32 in Internet. We recommend starting with [1]. Note that FAT32 terminology is using the term “cluster” instead of the term “block”. A typical cluster size is 8 sectors, that means 4 KB. In this case, a cluster will occupy 8 consecutive sectors. A sector is 512 bytes long. A disk that has a file system on it is called a *volume*.

In this project you will work with a FAT32 volume that will not sit on a separate physical media but on a regular Linux file. Therefore, you first need to create a regular Linux binary file that will act as your disk, i.e. a virtual disk, with Linux dd command. The dd command will initialize the file to all zeros. Block size will be 4KB. You can create such a file with as many blocks as you wish, of course there is a limit, but in this example we will create a file that has 1024 blocks. The name of the virtual disk will be “vdisk”.

**dd if=/dev/zero of=vdisk bs=4k count=1024**

Linux file vdisk will act as a disk: a virtual or pseudo disk. We can format it with FAT32 file system and mount it to an empty directory in our local Linux file system.

Now use the **losetup** command to associate a loop device file (/dev/loop0) with the virtual disk file. In this way you make the virtual disk file look like a block storage device instead of just a regular binary file within the file system [5].

**losetup /dev/loop0 vdisk**

Each device in a Linux computer has a device file associated with it. A device file starts with the /dev prefix. After running losetup command as above, the virtual disk file is now treated as a block oriented storage device that can be mounted with the mount command. To delete the association between the loop device file and the virtual disk file, we can use the -d option of losetup as: **losetup -d vdisk**.

You can now create a FAT32 file system in your vdisk file. This is called formatting the disk. You can use the mkfs.fat command of Linux for this purpose, which may take several options. Run the command with options given in the example below. We will always use the same options in our tests. You are suggested to read the man page of mkfs.fat or mkdosfs commands.

**mkfs.fat -a -v -S 512 -s 8 -F 32 -n CS342 vdisk**

Above, the `-F` option says that the file system will be FAT32, not FAT16 or FAT12, `-s 8` options says that the block size will be 8 sectors, that means 4KB, `-S 512` says that a logical sector is 512 bytes long, and `-n` option gives a name for your volume. Hence after creating the file system, you now have a FAT32 volume which you can mount and use; or you can write a program to analyze the volume directly without mounting.

Create a mount point, that means an empty directory with the `mkdir` command. The name of the empty directory can be, for example, “myfat32fs”. Then mount your volume on this mount point. Read the man page of the mount command.

```
sudo mount -t msdos ./vdisk myfat32fs
```

You can now change into myfat32fs directory. That means you are changing into the root directory of the FAT32 volume you mounted. There you can create and use files with any tools that you like: touch command, dd command, cp command, an editor, or something else. You can also create subdirectories.

Now, unmount your volume.

```
sudo umount vdisk
```

You can now analyze the virtual disk file. You will write a program, called `fat32.c`, that will read and manipulate it directly one byte at a time, or 512 bytes at a time, that means one sector at a time, or 4096 bytes at a time, that means one block at a time. The executable will be called **fat32**.

As said, your `fat32` program will read the virtual disk file block by block, or sector by sector and will do some operations on the file system structures directly. That means you will be accessing your virtual disk in raw mode in your program. In our test cases, we will have just have one directory, i.e., a root directory, in a FAT32 virtual disk. There will be no subdirectories and the number of files in the root directory will not exceed 50. Hence the root directory will be fitting into a single block. You can assume that only *short filenames* will be used. A short filename has at most 8 characters for the filename and 3 characters for its extension. “project.pdf”, for example, is a short filename. We will not use long filenames.

The program will be invoked in one of the following ways, doing the specified things:

- **fat32 <vdiskfile> -p volumeinfo**: Print information from the first sector (VolumeID sector) of the FAT32 file system. Printed information will include some numbers, like the number of sectors that FAT occupies.
- **fat32 <vdiskfile> -p rootdir**: Print the content of the root directory. For each directory record you will print the filename, its extension, and also other information found in the record, as nicely as possible. You will use one output line per record.
- **fat32 <vdiskfile> -p blocks <filename>**: Print the numbers of the clusters, i.e., blocks allocated to file <filename>. If file size is 0, nothing will be printed. One number per line will be printed.
- **fat32 <vdiskfile> -d <filename>**: Delete the file with name <filename>. The file to delete is a FAT32 file in the volume. You will search the root directory to find the related information. Your delete operation will delete the corresponding directory record, that means will mark it as deleted, and will update the FAT so that the corresponding blocks appear as free. Not that there are two FAT tables in a volume for reliability reasons.

Sample invocations of your program can be as follows:

- **fat32 vdisk -p volumeinfo** : Print info from the volumeID sector.

- **fat32** vdisk -p rootdir : Print info about root directory content.
- **fat32** vdisk -p blocks x.bin : Print numbers of the blocks of the file with name "x.bin".
- **fat32** vdisk -d x.bin : Delete the file "x.bin".

You will be provided an example virtual disk, called vdisk, in the webpage of the course. It is a regular Linux binary file. It acts as a FAT32 volume of size 4MB. That is like a 4MB disk formatted with FAT32. In that we have some files created and sitting. As an exercise you can start with parsing this file. You can see the content of the vdisk file in hex using the command xxd.

```
xxd vdisk
```

### Submission:

Submit through Moodle. You will do a demo and you will be asked very detailed questions.

### Clarifications:

- May be added to webpage.

### References:

- [1]. Understanding FAT32 File Systems.  
<http://www.pjrc.com/tech/8051/ide/fat32.html>.
- [2.] FAT32 File Systems Specification.  
<https://msdn.microsoft.com/en-us/library/gg463080.aspx>
- [3]. Design of FAT Filesystems.  
[https://en.wikipedia.org/wiki/Design\\_of\\_the\\_FAT\\_file\\_system](https://en.wikipedia.org/wiki/Design_of_the_FAT_file_system)
- [4]. FAT file systems: <http://wiki.osdev.org/FAT>
- [5]. Anatomy of the Linux File system.  
<http://www.ibm.com/developerworks/linux/library/l-linux-filesystem/>